

Everything that runs, what it costs, and what breaks.

A full technical account of the platform as it stands on 17 August 2026 — the week's changes and the reasoning behind them, the unit economics measured against live provider prices, the metrics worth watching, and the sequenced work that follows.

Headline: the code is live and both media paths are now proven.

Commit `e831e05` is on `main`, pushed, and auto-deployed by Hostinger — production is serving this week's work. Every required readiness check is green, **including private media storage, which was the blocker**. A video has since been generated, saved and played back from a real account, and the image path has been verified end to end with a real 1.7MB generated file. The central feature of the product is repaired in fact, not in theory.

PREPARED

17 August 2026

COMMIT UNDER REVIEW

e831e05 (deployed)

ENVIRONMENT

ai360.africa · Hostinger · Supabase

READINESS

200 · all required checks green

01

Where the system stands

Every figure below is read from production — the database, the live readiness endpoint, the deployed HTML and the provider's own catalogue. Nothing is estimated.

11

accounts, 9 workspaces. None created since 14 August.

GH¢125

collected. One Everyday subscription active.

0

images or videos ever delivered to a person.

Platform state, verified

COMPONENT	STATE	EVIDENCE
Deployment	Live	e831e05 on origin/main; live HTML serves the new nav, hero and CTAs
Readiness endpoint	200 ready	app_url, ai_gateway, supabase_auth, database, media_storage, payments all green
Private media storage	Configured	New media_storage check reports ready from production
Media actually delivered	Proven	Video generated, stored and played back; image path verified with a real 1.7MB file
Error tracking	Live	Sentry client active on ai360.africa with a real DSN
Searchable logs (Axiom)	Not provisioned	No dataset key in the production environment
Transactional email	Off	EMAIL_ENABLED=false — paying customers get no receipt
Payments	Open	ExpressPay live; 1 approved GH¢125, 1 pending
Agent runs	Zero	lab_agent_runs = 0; the agent pillar has never been used

The media ledger, in full

Seven media jobs have ever existed. Every one of them failed to deliver, and the reason was recorded by production itself.

TYPE	COUNT	STATUS	RECORDED REASON
Image	4	failed	Private media storage is not configured.
Video	2	completed, no file	Settled at 19 credits each with no stored output
Video	1	running	delivery_retry — clip exists at the provider, recoverable

Consequence to settle: 38 credits were charged for two clips that were never delivered. They are owed either the finished files or a refund. The third clip is still downloadable from the provider and can be delivered now that storage works.

02 What changed this week, and why

Six decisions and one incident, each recorded in `DECISIONS.md` with its reasoning. Presented as before and after, because in every case the old behaviour was defensible-looking and wrong.

1 · Private media storage was never configured in production

WAS

Renders completed at the provider and were then lost. The studio showed “Rendering...” forever because a failed save was treated as a passing fault and retried indefinitely. Images appeared to work because they fall back to a temporary base64 copy that is never stored. The real error existed only in host logs.

NOW

`MediaStorageNotConfiguredError` is typed and terminal: the job fails, the full credit hold is refunded, and the person is told. Delivery that is still failing 30 minutes after the provider finished is also terminal. The real reason is written to the job row. `/api/ready` gained a required `media_storage` check.

2 · Video had one price when it should have had two

WAS

Every render silently used the expensive engine at 19 credits. A cheaper engine existed in the catalogue and was never offered. The credit floor of 12 meant even a cheap clip was charged nearly the full rate, so choosing it would have bought the customer little.

NOW

The quote drives the hold. Floor cut 12 → 6, ceiling raised 20 → 24. A draft clip costs about 7 credits against 19 for the best engine — roughly a third — and the ceiling change lifts engine headroom from 7% to 23%.

3 · The landing page promised what a visitor could not have

WAS

Three of four hero starters and four of six outcome cards routed to `agent` or `studio`, which require an account. A signed-out visitor clicking “Take the first step” received `POST /api/agent → 401` and an error card with a UUID as their first experience. The page said “Start free, no card”.

NOW

Free chat paths lead; gated starters carry a “Free account” badge before they are clicked. An empty click runs a real question on the free path instead of opening another empty box. Pre-written prompt stems no longer submit as goals.

4 · Eight labels pointed at one destination

WAS

“Take the first step”, “Start with your goal”, “Start with a goal”, “Start now”, “Try it now”, “Open AI360”, “Open workspace”, “Continue as a guest” — all to `/app`, each written by hand. The hero submitted immediately while outcome cards only prefilled, with no rule distinguishing them.

NOW

One `StartCta` owns label, destination and send-or-prefill. “Start free” signed out, “Open workspace” signed in. The rule is explicit: what the person typed runs; what we pre-wrote waits.

5 · Naming and pilot language

WAS

Nav said “Plans”, the URL was `/pricing`, metadata said both. Paid surfaces described a private trial — “One-month pilot access”, “Join the Everyday pilot” — while charging real money. The terms claimed prices were “a pilot proposal until checkout is activated” and that top-ups were not sold. Both were untrue.

NOW

The page is Pricing; the products on it are plans. Pilot language is gone from every charged surface while the real constraints stay stated: one month per payment, no automatic renewal, no annual billing yet.

6 · The type scale was a quarter too small

WAS

Body copy 15px. On the landing page at 1280px, 76 of 166 text elements sat below 12px, the smallest at 7px. Across all stylesheets, 432 of 883 `font-size` declarations were under 12px. The product looked better at 125% browser zoom because zoom was correcting a real defect.

NOW

Every `font-size` scaled on a curve, heaviest at the small end, 12px floor. Body is 17px; nothing is under 12px; longest line fell from 90 to 72 characters. Spacing untouched, so no layout could reflow. Verified at 375, 768, 1024, 1280 and 1440px.

A seventh change nobody asked for: cost telemetry

The video status endpoint recorded the provider's full job cost on every poll. 150 polls of three renders were recorded as **\$47.80 of spend against roughly \$0.96 of real cost** — a fifty-fold overstatement, in exactly the number a pricing decision would be made from. Now only a terminal status carries the cost, and only once.

03

Unit economics, measured

Costs below come from the live OpenRouter catalogue and from measured production spend, converted through the same module that charges customers. One credit is defined as GH¢0.26 of landed cost at \$1 = GH¢13, including the platform fee and a 10% currency buffer.

What a credit sells for

PACK	PRICE	CREDITS	PER CREDIT	GROSS MARGIN
Everyday	GH¢125	120	GH¢1.04	75%
Builder	GH¢350	400	GH¢0.88	70%
Team	GH¢1,200	1,400	GH¢0.86	70%
Top-up 50	GH¢50	40	GH¢1.25	79%
Top-up 100	GH¢100	90	GH¢1.11	77%
Top-up 200	GH¢200	185	GH¢1.08	76%

Every pack clears cost several times over, and top-ups correctly stay dearer per credit than a subscription, so nobody is better off avoiding a plan. A fully-consumed Everyday subscription costs GH¢31.20 of model spend against GH¢125 collected.

Cost per unit of work, against the live catalogue

WORK	ENGINE	PROVIDER COST	CREDITS OF COST	CHARGED	POSITION
Video, 4s Quick	veo-3.1-lite	\$0.1200	7	7 credits	shown before choosing
Video, 4s Better	veo-3.1-fast	\$0.3200	19	19 credits	shown before choosing
Video, 4s Best	veo-3.1	\$0.8000	47	47 credits	now sellable
Video, 8s Quick	veo-3.1-lite	\$0.2400	14	14 credits	cheaper than a 4s Better
Video, 8s Best	veo-3.1	\$1.6000	93	93 credits	inside the 96 per-second cap
Image	gpt-image-1-mini	\$0.0333	2	3 credits (floor)	measured, 85% margin
Chat turn	fast default	\$0.0012	<1	included	bounded by daily cap
Agent run	coordinator	\$0.0089	1	3–8 credits	wide margin

The ceiling is now per second, not per render. A flat cap could not describe an engine that charges by the second: at 24 credits it refused an eight-second clip outright, and it also made the premium engine unsellable, so “Best” silently resolved to the same engine as “Better” — a tier that existed in name only. At 12 credits a second the four-second cap is unchanged at 48,

every length from 4 to 8 seconds is priced proportionally, and the promise that nothing costs more than the person saw is carried by the quote they accept rather than by the cap. Re-run `scripts/audit-media-pricing.mjs` whenever the provider moves prices.

Real spend to date

Corrected for the polling over-count, total provider spend across the entire life of the product is approximately **\$1.40** — about GH¢21. The recorded figure of \$48.22 was an artefact. This matters because it is the denominator for every efficiency claim.

04

Technical decisions and the guardrails they set

Each of these is now written into `DECISIONS.md` as a rule, not a fix, because the same class of failure recurred three times this week.

Storage is a readiness requirement

If a deliverable only exists once it is stored, storage is a required health check — not an assumption. Production reported "ready" for two days while unable to keep a single file.

A retry loop must classify its faults

Transient and permanent failures need different handling. An unset environment variable cannot be retried into working, and treating it as transient produced an infinite spinner.

No surface holds its own copy of a price

Four separate drifts this week: the pricing page, the assistant's product knowledge, the chat daily caps, and the credit guide. Every published number is now derived from the module that charges.

A call to action is a promise about the next screen

"Take the first step" is only honest when a first step follows. Gating and copy are one decision, never two.

Typed errors over string matching

`MediaStorageNotConfiguredError` lets callers branch on cause rather than parsing messages, which is what allowed a permanent fault to be classified correctly.

Settlement must never break delivery

An unguarded `settleReservation` in the status path could turn every later poll into a 500 the client retried forever — so a delivered clip could still read as "Rendering...".

Diagnostic tooling added

The storage incident took a database query, a provider replay and a bucket write probe to isolate. That work is now one command each, so the next occurrence is minutes rather than hours.

SCRIPT	ANSWERS
<code>diagnose-video-jobs.mjs</code>	Durable jobs, credit holds, live provider status, and a bucket write probe
<code>diagnose-video-delivery.mjs</code>	Replays download → upload → rows against a stuck job; rolls back by default, <code>--commit</code> to deliver
<code>diagnose-media-storage.mjs</code>	What production itself recorded when storage failed
<code>verify-media-roundtrip.mjs</code>	Proves image and video can be stored and read back, then cleans up
<code>audit-media-pricing.mjs</code>	Live provider prices against credit charges and ceiling headroom
<code>report-status-metrics.mjs</code>	Aggregate-only production metrics for reporting

05

The measurement system

Four pillars. For each: what to watch, where the number actually lives, where it stands today, and the threshold that should trigger action. Anything without a data source is listed honestly as a gap rather than a metric.

Pillar 1 • Unit economics per feature

Metric	Source	Today	Threshold
Gross margin per credit, by pack	audit-media-pricing.mjs	70–79%	Act below 60%
Engine headroom under ceiling	Live video catalogue	23% standard	Act below 15%
Measured cost vs credits charged, per feature	lab_usage_events	Image 87% margin	Act if charged < cost
Cost per active account per month	usage events ÷ accounts	~\$0.13	Act above 25% of revenue
Provider price drift since last audit	OpenRouter catalogue	Audited today	Re-audit monthly

Pillar 2 • Reliability and incidents

Metric	Source	Today	Threshold
Media delivery success rate	lab_media_jobs vs lab_media_outputs	0 of 7	Below 95% is an incident
Jobs terminal without an output	lab_media_jobs	3	Any occurrence is investigated
Charged-without-delivery events	Reservations vs outputs	2 (38 credits owed)	Must be zero
Silent-failure classes open	Manual register	1 (image data-URL fallback)	Each needs a readiness check
Time to detect	Sentry vs incident log	~2 days for storage	Target under 1 hour
Readiness endpoint status	/api/ready	200	Any 503 pages someone

The lesson to instrument: the storage failure was invisible for two days because nothing watched the gap between "job completed" and "file exists". That comparison is now the single most valuable query in the system and belongs on a daily cron.

Pillar 3 • Activation and retention

STAGE	DEFINITION	TODAY	NOTES
Visit → first message	Any chat turn sent	15 conversations / 11 accounts	Reasonable early signal
First message → real output	Export, image or video kept	0	The funnel's actual hole
Return visit within 7 days	Second session	Not instrumented	gap
Free → paid	Approved payment	1 of 11 (9%)	One controlled live test
Agent adoption	Any agent run	0 of 11	A whole pillar unused
New accounts per week	Sign-ups	0 since 14 Aug	No acquisition motion yet

Pillar 4 · Roadmap progress

PHASE	OUTCOME	RECORDED STATUS	REALITY CHECK
0	Decisions and baseline	Complete	Holds
1	Identity and tenant isolation	Blocked on credentials	Stale — Supabase Auth shipped; workspace isolation live
2	Durable projects, assets, usage	In progress	Assets path now provably works but has never carried real output
3	Coordinator and specialists	Planned	0 agent runs — unvalidated with users
4	Live agent-room experience	Planned	Unstarted
5	Evals, security, cost controls	Planned	Partly delivered ahead of plan: Sentry, cost audit, readiness checks
6	Rollout and subscriptions	Planned	Overtaken by events — checkout is open and money has moved

The roadmap has drifted from reality in both directions: Phase 1 is further along than recorded, Phase 6 has partly happened out of order, and Phase 5 work landed early because incidents forced it. It needs a rewrite against what is actually true.

06

Risk register

RISK	SEVERITY	WHY IT MATTERS	MITIGATION
Media delivery unproven in production	High	The paid, differentiating feature has never delivered a file to anyone	Generate one image and one video today and confirm both persist
Paying customer receives no receipt	High	<code>EMAIL_ENABLED=false</code> while real money is being taken	Complete domain verification and enable
Image cost is measured on one model only	Low	The primary engine is verified working at \$0.0333 an image, but the fallback's real cost has not been measured separately	Record both in the pricing audit
Single viable video engine	Medium	Standard tier has 23% headroom; premium is unsellable at current ceilings	Monthly pricing audit; consider a second provider
Agent pillar unvalidated	Medium	Zero runs. A core promise of the product is untested by users	Guided agent task in onboarding
Duplicate page render	Medium	Every marketing page ships a hidden second copy — ~17.7KB of 66KB, duplicate DOM ids — on the slow connections we target	Own investigation; streaming-level cause
No acquisition motion	Medium	Flat sign-ups since 14 August; all accounts from our own circle	Decide the first channel deliberately
Type scale unreviewed by eye	Low	Verified by computed geometry across five widths, but not visually confirmed	Ten-minute human pass at 100% zoom

07 Next steps, sequenced

Ordered by what unblocks the most. The first block is measured in minutes and settles the question the whole product rests on.

Today · prove the repair

1. Generate one image in production. Confirm a row appears in `lab_assets` and the file opens from the gallery.
2. Generate one 4-second draft video. Confirm it persists, the credits settle to the measured cost, and the clip plays.
3. Deliver the stuck clip with `node scripts/diagnose-video-delivery.mjs --commit` then settle the 38 credits owed on the two undelivered renders.
4. Look at the new type scale at 100% zoom on a phone and a laptop, and confirm it matches what you saw at 125%.

This week

1. **Turn on transactional email.** Complete the domain verification step and set `EMAIL_ENABLED=true`, so a paying customer gets a receipt without one being sent by hand.
2. **Generate one image through the signed-in studio.** The provider, guards, upload, asset row and read-back are each verified with a real 1.7MB image; the only untested link is the authenticated route wrapper, which video has already exercised live.
3. **Add the delivery-gap query to a daily cron.** Terminal jobs without a stored output, and reservations settled without delivery. This is the check that would have caught the storage failure in an hour rather than two days.
4. **Provision the Axiom dataset** so logs are searchable rather than only alertable through Sentry.
5. **Rewrite the roadmap** against what is actually true, with the phase drift resolved.

Next two to four weeks

1. **Carry the goal through sign-up.** A visitor who picks a gated starter should sign up and have their goal run immediately, rather than losing it at the wall.
2. **Resolve the duplicate render.** Roughly a quarter of every marketing page's HTML is a hidden second copy, with duplicate DOM ids, on connections where weight is expensive.
3. **The spacing pass.** The other half of the 125% effect. Higher layout risk than the type change, so it needs the same width-by-width verification.
4. **Instrument return visits** so retention becomes measurable rather than assumed.
5. **Get the first agent run from a real user** — offer a guided agent task in onboarding and watch what happens.

Beyond, gated on evidence

1. **Annual billing**, only after repeat payments, reversals and refunds are proven end to end.
2. **Second media provider** for engine redundancy and pricing leverage.
3. **Team workspaces**, currently flag-disabled, once single-user retention is real.
4. **The moat work:** local context, language coverage and product self-knowledge are the parts a global competitor cannot copy cheaply.

This codebase — formerly “AI360 Lab”, now simply AI360 — carries four institutional jobs at once. They are compatible, but they are measured differently, and confusing them is how a product ends up optimising for the wrong thing.

1 • A commercial product Revenue from Ghanaian individuals and businesses, at 70–79% gross margin per credit. Proven once, with GH¢125 collected end to end. The economics work; the volume does not exist yet. Measured by: paying customers, margin per credit, cost per active account.	2 • Digital inclusion and capability Frontier AI made usable for students, traders, NGOs and public servants — in plain language, on a phone, on an interrupted connection, priced in cedis with Mobile Money. Most of the product’s hardest engineering serves this rather than any feature list. Measured by: completed real outcomes, not sign-ups.
3 • Proof of AIC’s capability A live production system with credit metering, payment reconciliation, durable job recovery, error tracking and an incident discipline. This week it detected, diagnosed and fixed a silent revenue-affecting fault from its own records. That is the evidence that AIC can operate software, not only build it. Measured by: uptime, incident quality, decision record.	4 • An investment and partnership case The architecture reads as fundable: honest unit economics, enforced spend ceilings, no silent overage, provider independence, and a written record of every trade-off. What is missing is traction, and traction is now the binding constraint rather than capability. Measured by: growth rate, retention, unit economics at scale.

The strategic read

The build risk is largely behind us; the demand risk is entirely ahead.

The platform can take money, meter work, recover from failure and tell the truth about cost. It has eleven accounts, one payment and no delivered creative output. Every remaining question is about whether people in Accra will use and pay for this — and none of them can be answered by more engineering.

Which argues for a deliberate order: prove media delivery today, put the product in front of a first cohort of twenty real users this month, and let what they actually do decide the next phase of the roadmap rather than the plan written before they existed.